

HW3.1. Floating Point Theory

Feel free to check out the [guide](#) that we have prepared to help you in this problem.

For the following question, we will be following the IEEE-754 Floating Point standard *but with different numbers of bits*: with 6 bits of exponent, an exponent bias of -31, and 25 bits of mantissa. Due to PrairieLearn restrictions, please submit your answer as the full number.

Floating point representations present a trade-off between range and accuracy. Due to the limited significant digits that you can represent, there will be a point where integers cannot be exactly represented anymore.

Q1.1: What's the smallest non-infinite positive integer (whole number) this representation CANNOT represent?

*Hint: Think about how your floating point value changes when you change the mantissa by 1.*

integer

?

The floating point standard optimizes the representation of numbers by not including the implied 1 for the 'normal' binary representations.

Q1.2: What power of 2 is the smallest representable positive normalized number? Submit the exponent only.

integer

?

Floating point also allows for representations of numbers even smaller than the smallest normalized number. Denormal numbers utilize an implied 0, instead of an implied 1.

Q1.3: What power of 2 is the smallest representable positive value? Submit the exponent only.

integer

?

Floating point also has support for special values that can be returned as the result of invalid operations.

Q1.4: How many NaNs can be represented with this floating point system?

*Hint: Recall how NaNs are represented following the floating point standard.*

integer

?

Q1.5: What's the most negative possible denormalized number representable in this system? Write the value in IEEE-754 hexadecimal notation (with uppercase letters), including the "0x" prefix.

?

The next few questions explore some quirkier aspects of floating-point numbers.

Q2.1: Addition and subtraction on unsigned and two's complement numbers is both commutative and associative. (You can arbitrarily re-order a series of additions and subtractions and get the same result.) In contrast, floating point addition is **not** associative, because between each substep, the result loses a small amount of precision. Given the floating point encoding from earlier, which of the following sequence of floating point additions and subtractions returns the correct result?

- ☐ (a)  $(2^{-10} + 2^{22}) + (-2^{22} + 2^{-1})$
- ☐ (b)  $2^{-10} + (2^{22} + (-2^{22} + 2^{-1}))$
- ☐ (c)  $2^{-10} + 2^{22} - 2^{22} + 2^{-1}$

□

Homework 3

Assessment  
overview

Total points: 0/100  
Score: 0%

Question

Value: 50

History:

Awarded points: 0/50

Report an error in this question

Previous question

Next question

Attached files

No attached files

Attach a file

Attach text

⌋ (d)  $2^{-10} + (2^{22} - 2^{22}) + 2^{-1}$

Select all possible options that apply. ?

Q3.1: An IEEE-754 **single-precision** floating point number is 32 bits, consisting of 1 sign bit, 8 exponent bits, and 23 mantissa bits. Can this floating-point encoding represent the space of all integers a 32-bit, two's complement number can represent?

- ☐ (a) No
- ☐ (b) Yes

Save & Grade 20 attempts left

Save only

Additional attempts available with new variants ?